

An AI brand-intelligence analyst that shows its sources

Every agent answer, traceable to a source. Agnel Salve · Northeastern University · Fall 2026

Problem

Weave's brand and data teams track reputation and competitors by hand across news, analyst notes, legal notices and practitioner chatter — slow, and easy to miss the story that matters. AI summaries could do it faster, but nobody can trust them, because nothing ties a claim back to a source.

Solution approach

- **Collect** ~500 items a run from 17 public feeds — Weave, 5 competitors, AI-market news, research, Reddit, Hacker News (A3 pipeline, extended).
- **Judge** every item with Gemini 3.5 Flash-Lite: relevance, sentiment, risk, opportunity, brand archetype, key fact.
- **Decide & act:** high risk → instant Gmail alert; competitor stumble → opportunity alert; repeats suppressed.
- **Write** cited briefs with Gemini 3.8 Flash — numbers computed in code, never generated.
- **Fact-check** every claim with a *different* model; unsupported claims are removed and the brief is scored.
- **Survive failure:** retries, model chain on quota, JSON repair, circuit breaker, budget cap, error workflow.

Business value

- Replaces an estimated **4–6 analyst hours a week** of reading, writing and source-checking for about **\$0.16 of AI per run**.
- Risks reach the inbox the run they appear — e.g. the shareholder-investigation alert on the \$650M take-private.
- Every sentence links to its source and carries a groundedness score, so the team can act on it.
- At 1,000 API requests/day: **≈ \$9–11 a month** in AI (measured per-request cost).

Current performance (measured 20261009)

294

RECORDS JUDGED
PER RUN

**301s →
86s**

COLD → CACHED
RUN

\$0.16

AI COST PER RUN AT
LIST PRICE (\$0 ON
FREE TIER)

87%

MEDIAN
GROUNDEDNESS, 12
BRIEFS FACT-
CHECKED (54–97%)

13

NOISE ITEMS THE AI
FILTERED OUT

6

REAL EMAILS
DELIVERED TO
GMAIL (3 OF 3
TYPES)

100%

SUCCESS AT 100
CONCURRENT AI
REQUESTS

1,000

REQUEST BURST
N8N HANDLES
(100%)

Sample outputs

